



Parallelization of a digit recognition program

---Using python MPI---

Course : Parallel and distributed computing

Mahima M Rao

6th Sem

Pes2201800257

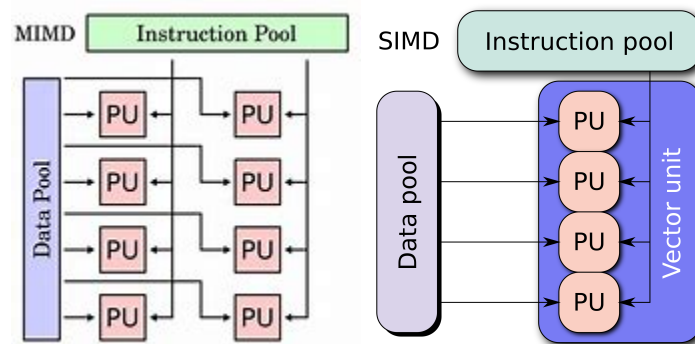
Pesu ECE

MPI - Message passing interface


→ MPI for Python provides an object oriented approach to message passing which grounds on the standard MPI-2 C++ bindings. The interface was designed with focus in translating MPI syntax and semantics of standard MPI-2 bindings for C++ to Python.

→ Some basics of MPI

- ◆ Commutators ---> Initiates communication between Process (nodes,host,computer cluster) Via message passing between the nodes .
- ◆ Works for both SIMD and MIMD






- ◆ Point to point communication : -
 - Using send and receive objects
 - Can be blocking and non blocking depends on object




Collective Communication

- ◆ Barrier synchronization across all group members.
- ◆ Global communication functions
 - Broadcast data from one member to all members of a group.
 - Gather data from all members to one member of a group.
 - Scatter data from one member to all members of a group.
- ◆ Global reduction operations such as sum, maximum, minimum, etc
- ◆ All to all , many to many , all to one , allgather , allreduce etc





MPI_Reduce(sbuf,rbuf,6,MPI_INT,MPI_SUM,0,MPI_COMM_WORL

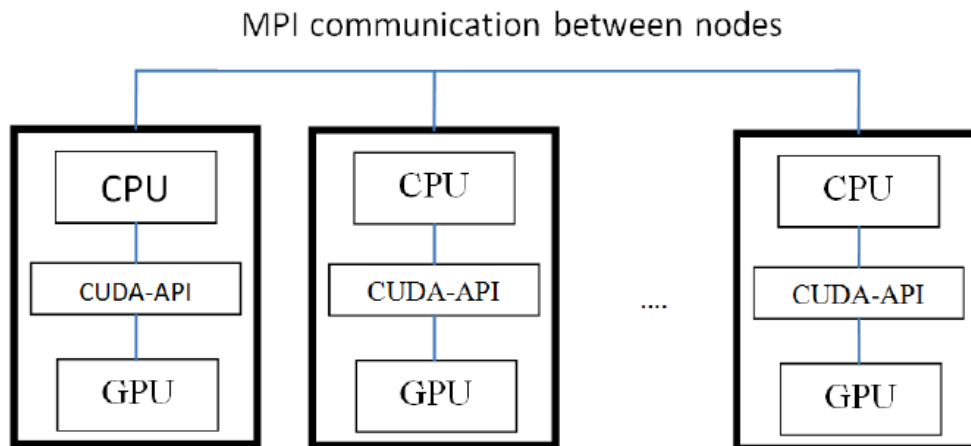


Novel paradigms are becoming more and more common:


MPI+OpenMP

MPI+GPU, etc --- This can be achieved with the CUDA support for python using “cupy” class

Where in it sends data to gpu calling out threads




Digit recognition using perceptron learning algorithm




→ This is the seven segment display

Consider the fully connected network



First - training all 10 nodes parallelly



Code:

```
import mpi4py
from mpi4py import MPI
import time
```

```
digit_R=[[0,0,0,0,0,0,0],
          [1,1,0,0,0,0,0],
          [1,0,1,1,0,1,1],
          [1,1,1,0,0,1,1],
          [0,1,0,0,1,0,1],
          [0,1,1,0,1,1,1],
          [0,1,1,1,1,1,1],
```

```
[1,1,0,0,0,0,0],  
[1,1,1,1,1,1,1],  
[1,1,1,0,1,1,1]]
```

```
y = [[1,0,0,0,0,0,0,0,0],  
      [0,1,0,0,0,0,0,0,0],  
      [0,0,1,0,0,0,0,0,0],  
      [0,0,0,1,0,0,0,0,0],          % expected output  
      [0,0,0,0,1,0,0,0,0],  
      [0,0,0,0,0,1,0,0,0],  
      [0,0,0,0,0,0,1,0,0],  
      [0,0,0,0,0,0,0,1,0],  
      [0,0,0,0,0,0,0,0,1],  
      [0,0,0,0,0,0,0,0,1]]
```

```
# Make a prediction with weights
```

```
def predict(row,weights):  
    activation = weights[0]  
    #print(len(row))  
    for i in range(len(row)):  
        activation += weights[i + 1] * row[i]  
    return 1.0 if activation >= 0.0 else 0.0
```

```
# Estimate Perceptron weights using stochastic gradient descent
```

```
def train_weights(train,y, l_rate, n_epoch):  
    weights = [0.0 for i in range(len(train[0])+1)]  
    cost=[]  
    for epoch in range(n_epoch):  
        sum_error = 0.0  
        for row,i in zip(train,y):  
            prediction = predict(row, weights)  
            error = i - prediction  
            sum_error += error**2  
            weights[0] = weights[0] + l_rate * error #bias update  
            for i in range(len(row)):  
                weights[i + 1] = weights[i+1] + l_rate * error * row[i]  
#weight update  
        #cost.extend([sum_error])  
    return weights
```

```

def perceptron_learning_algorithm(logic,y,name) :
    l_rate = 0.1
    n_epoch = 30
    print("-----%s-----" %(name))
    weights = train_weights(logic,y,l_rate, n_epoch)
    print("Trained weight = %s" %(weights))
    for row,i in zip(logic,y):
        prediction = predict(row, weights)
        print("Expected=%d, Predicted=%d" % (i, prediction))
    if rank == 0 :
        stop_time = time.time()
        print("Time consumed = ",stop_time - start_time)
    return(weights)

if __name__ == "__main__":

    comm = MPI.COMM_WORLD
    rank = comm.Get_rank()
    size = comm.Get_size()                % Parallelized part

    start_time = time.time()
    print(rank,size)

    r = perceptron_learning_algorithm(digit_R,y[rank],rank)
    a = comm.gather(r,root=0)

    print(a)

```

Results :

(mpi) C:\Users\mahim\Desktop\k mean>mpiexec -np 10 python mip_prog.py

Here i have used 10 process to parallely execute for training 10 different set of inputs
Each process are run independently

7 10

-----7-----

Trained weight = [-0.1, 0.1, 0.1, -0.1, -0.1, -0.1, -0.1, -0.1]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

Expected=0, Predicted=0

Expected=0, Predicted=0

None % here 7 is expected 7 is predicted and this is executed by 7th rank process

Similarly others

5 10

-----5-----

Trained weight = [-0.2, -0.6, -0.1, 0.2, -0.6, 0.0, 0.2, 0.0]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

None

3 10

-----3-----

Trained weight = [-0.20000000000000004, 0.1, 0.0, 0.1, -0.1, -0.2, 0.1, 0.0]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

None

2 10

-----2-----

Trained weight = [-0.1, 0.0, -0.1, 0.0, 0.1, 0.0, 0.0, 0.0]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

None

6 10

-----6-----

Trained weight = [-0.30000000000000004, -0.7999999999999999, -0.2, 0.1, 0.6, -0.1, 0.1, -0.1]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

None

1 10

-----1-----

Trained weight = [-0.2, -0.1, 0.0, -0.1, -0.1, 0.0, -0.1, -0.1]

Expected=0, Predicted=0

Expected=1, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

None

4 10

-----4-----

Trained weight = [-0.1, -0.1, 0.0, -0.2, 0.0, 0.1, -0.2, 0.1]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

None

9 10

-----9-----

Trained weight = [-0.5, 0.20000000000000004, -0.1, 0.0, -0.2, 0.6, 0.0, -0.2]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

None

8 10

-----8-----

Trained weight = [-0.6, 0.20000000000000004, 0.10000000000000003,
-0.20000000000000004, 0.6, 0.4, -0.20000000000000004, -0.20000000000000004]

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=1, Predicted=1

Expected=0, Predicted=0

None

0 10

-----0-----

Trained weight = [0.0, -0.1, -0.1, 0.0, 0.0, 0.0, 0.0, 0.0]

Expected=1, Predicted=1

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Expected=0, Predicted=0

Time consumed = 0.0065081119537353516

[[0.0, -0.1, -0.1, 0.0, 0.0, 0.0, 0.0, 0.0], [-0.2, -0.1, 0.0, -0.1, -0.1, 0.0, -0.1, -0.1], [-0.1, 0.0, -0.1, 0.0, 0.1, 0.0, 0.0, 0.0], [-0.20000000000000004, 0.1, 0.0, 0.1, -0.1, -0.2, 0.1, 0.0], [-0.1, -0.1, 0.0, -0.2, 0.0, 0.1, -0.2, 0.1], [-0.2, -0.6, -0.1, 0.2, -0.6, 0.0, 0.2, 0.0], [-0.30000000000000004, -0.7999999999999999, -0.2, 0.1, 0.6, -0.1, 0.1, -0.1], [-0.1, 0.1, 0.1, -0.1, -0.1, -0.1, -0.1, -0.1], [-0.6, 0.20000000000000004, 0.10000000000000003, -0.20000000000000004, 0.6, 0.4, -0.20000000000000004, -0.20000000000000004], [-0.5, 0.20000000000000004, -0.1, 0.0, -0.2, 0.6, 0.0, -0.2]]

Second - parallelizing the calculation of activation function

By performing :

$$Y = Y + X * W \text{ ---in all levels}$$

I,e parallelizing predict function as its accessed many times

Code :

% parallelized predict function

```
def predict(row,weights,rank,comm):
```

```
    a0 = weights[0]
```

```
    w = np.array(weights[1:len(row)+1])
```

```
    r = np.array(row[0:len(row)])
```

```
    a = w*r
```

```
    a1 = comm.allreduce(a,op=MPI.SUM)
```

```
    a = sum(a1) + a0
```

```
    return 1.0 if a >= 0.0 else 0.0
```

Results

Sequential : time consumed : 0.37s

First parallel : Ran it using 10 process : 0.0043 to 0.006s

Second Parallel : 0.34 s -- using 1 process --optimal , if used more than that due to extensive process communication runtime increases.

Combining both First and Second parallel : 1.12s again the same reason , run time decreases as number of nodes increases ---> **gives out good speed up for larger input vectors**